

# ISyE6420 Bayesian Statistics

## Project: Movie Recommendation System

Swati Gupta  
sgupta612  
GT ID : 903409986

April 22, 2019

### Introduction

The aim of the project is to build a movie recommendation system. Based on a given set of ratings for movies by users, we want to predict a user's rating for a movie and provide recommendations based on that. To do so, we will use Bayesian techniques and Gibbs sampling.

### Data

To build this model, we used a subset of the movie ratings data provided by GroupLens. The data was obtained from Kaggle. It contains close to 97000 ratings for over 6000 movies by more than 600 users. The data can be found here: [https://www.kaggle.com/rounakbanik/the-movies-dataset#ratings\\_small.csv](https://www.kaggle.com/rounakbanik/the-movies-dataset#ratings_small.csv). The data was cleaned to include movies which had at least 2 ratings and users who had rated at least 2 movies.

### Approach

The model we build will use collaborative filtering to predict the rating by a user for a movie. The principle behind collaborative filtering is that similar people will rate similar movies in a similar way. If A and B both like a movie X, and A also likes movie Y, B's likeliness towards Y would be closer to A's than any random user.

To evaluate similarity between users or movies, we introduce latent variables in our model. The latent variables will capture properties of a movie that makes a user like or dislike it. Some examples of such properties could be popularity of the cast, duration of the movie, amount of humor in the movie, etc.

# Defining the model

The following variables are used in the model

- $L$ : The number of latent factors
- $U$ : The number of users
- $M$ : The number of movies
- $N$ : The number of observed ratings
- $Y_{um}$ : The rating given to movie  $m$  by user  $u$
- $Y$ : The full vector of observed ratings
- $\gamma_m$ : A vector of length  $L + 1$  for movie  $m$ . The first element  $\gamma_m[0]$  is the bias for the movie. The remaining  $L$  elements,  $\gamma_m[1 :]$ , are the latent factors associated with the movie
- $\Gamma$ :  $M \times (L + 1)$  matrix where row  $m$  is  $\gamma_m$
- $\theta_u$ : A vector of length  $L + 1$  for user  $u$ . The first element  $\theta_u[0]$  is the bias for the user. The remaining  $L$  elements,  $\theta_u[1 :]$ , are user's preferences for the latent factors
- $\Theta$ :  $U \times (L + 1)$  matrix where row  $u$  is  $\theta_u$
- $\mu$ : The overall mean of ratings
- $\sigma^2$ : The residual variance of ratings

We define the following model to predict the ratings

$$\begin{aligned}
 Y_{um} &= \mu + \theta_u[0] + \gamma_m[0] + \theta_u[1 :]^T \gamma_m[1 :] + \epsilon_{um} \\
 \epsilon_{um} &\sim N(0, \sigma^2) \\
 p(Y_{um} | \mu, \sigma^2, \theta_u, \gamma_m) &= \frac{1}{\sqrt{2\pi\sigma^2}} \exp \left( - \frac{(y_{um} - (\mu + \theta_u[0] + \gamma_m[0] + \theta_u[1 :]^T \gamma_m[1 :]))^2}{2\sigma^2} \right)
 \end{aligned}$$

## Priors

The following priors are assumed for the parameters

$$\begin{aligned}
 p(\mu) &\propto 1 \\
 p(\sigma^2) &\propto \frac{1}{\sigma^2} \\
 p(\gamma_m) &= N(0, \Sigma_\gamma) \\
 p(\theta_u) &= N(0, \Sigma_\theta)
 \end{aligned}$$

where  $\Sigma_\gamma$ ,  $\Sigma_\theta$  are  $L + 1 \times L + 1$  covariance matrices. They are assumed to be diagonal matrices with all diagonal elements equal to  $\lambda_\gamma$  and  $\lambda_\theta$  respectively both assumed to be 1.

## Conditionals

$$p(Y, \mu, \sigma^2, \Theta, \Gamma) = \prod_{i=1}^N \frac{1}{\sqrt{2\pi\sigma^2}} \exp\left(-\frac{(y_i - (\mu + \theta_{u_i}[0] + \gamma_{m_i}[0] + \theta_{u_i}[1:]^T \gamma_{m_i}[1:]))^2}{2\sigma^2}\right) \\ \times \frac{1}{\sqrt{(2\pi)^n |\Sigma_\gamma|}} \exp\left(-\frac{1}{2} \gamma_m^T \Sigma_\gamma^{-1} \gamma_m\right) \frac{1}{\sqrt{(2\pi)^n |\Sigma_\theta|}} \exp\left(-\frac{1}{2} \theta_u^T \Sigma_\theta^{-1} \theta_u\right) \frac{1}{\sigma^2}$$

### Conditional for $\mu$

$$p(\mu|Y, \sigma^2, \Theta, \Gamma) = N(\mu_n, \sigma_n^2) \\ \text{where } \mu_n = \frac{1}{N} \sum_{i=1}^N (y_i - (\theta_{u_i}[0] + \gamma_{m_i}[0] + \theta_{u_i}[1:]^T \gamma_{m_i}[1:])) \\ \sigma_n^2 = \frac{\sigma^2}{N}$$

### Conditional for $\sigma^2$

$$p(\sigma^2|Y, \mu, \Theta, \Gamma) \propto (\sigma^2)^{-\frac{n}{2}-1} \exp\left(-\frac{1}{2\sigma^2} \sum_{i=1}^N (y_i - (\mu + \theta_{u_i}[0] + \gamma_{m_i}[0] + \theta_{u_i}[1:]^T \gamma_{m_i}[1:]))^2\right) \\ p(\sigma^2|Y, \mu, \Theta, \Gamma) = IG(\alpha_n, \beta_n) \\ \text{where } \alpha_n = \frac{N}{2} \\ \beta_n = \frac{\sum_{i=1}^N (y_i - (\mu + \theta_{u_i}[0] + \gamma_{m_i}[0] + \theta_{u_i}[1:]^T \gamma_{m_i}[1:]))^2}{2}$$

### Conditional for $\gamma_{m'}$

The term  $\gamma_{m'}$  appears only in the terms where a user rated the movie  $m'$  or in the prior itself.

$$Y_{um'} = \mu + \theta_u[0] + \gamma_{m'}[0] + \theta_u[1:]^T \gamma_{m'}[1:] + \epsilon_{um'}$$

We can rewrite the model as

$$Y_{um'} - \mu - \theta_u[0] = \gamma_{m'}[0] + \theta_u[1:]^T \gamma_{m'}[1:] + \epsilon_{um'}$$

Assuming we have  $g$  users who rated the movie  $m'$ , we can treat this as a linear regression

$$Y_{m'} = \gamma_{m'} X_{m'}$$

where

$$Y_{m'} = \begin{bmatrix} Y_{u_1 m'} - \mu - \theta_{u_1}[0] \\ Y_{u_2 m'} - \mu - \theta_{u_2}[0] \\ \vdots \\ Y_{u_g m'} - \mu - \theta_{u_g}[0] \end{bmatrix}, X_{m'} = \begin{bmatrix} 1, \theta_{u_1}[1:]^T \\ 1, \theta_{u_2}[1:]^T \\ \vdots \\ 1, \theta_{u_g}[1:]^T \end{bmatrix}$$

Taking into account the prior on  $\gamma_{m'}$ , we get

$$p(\gamma_{m'}|Y, \mu, \sigma^2, \Theta, \Gamma_{-m'}) = N(\mu_{m'}, \Sigma_{m'})$$

$$\text{where } \mu_{m'} = \Sigma_{m'} \frac{1}{\sigma^2} X_{m'}^T Y_{m'}$$

$$\Sigma_{m'}^{-1} = \frac{1}{\sigma^2} X_{m'}^T X_{m'} + \Sigma_{\gamma}^{-1}$$

### Conditional for $\theta_{u'}$

The term  $\theta_{u'}$  appears only in the terms where a movie is rated by user  $u'$  or in the prior itself.

$$Y_{u'm} = \mu + \theta_{u'}[0] + \gamma_m[0] + \theta_{u'}[1:]^T \gamma_m[1:] + \epsilon_{u'm}$$

We can rewrite the model as

$$Y_{u'm} - \mu - \gamma_m[0] = \theta_{u'}[0] + \theta_{u'}[1:]^T \gamma_m[1:] + \epsilon_{u'm}$$

Assuming we have  $g$  movies that were rated by the user  $u'$ , we can treat this as a linear regression

$$Y_{u'} = \theta_{u'} X_{u'}$$

where

$$Y_{u'} = \begin{bmatrix} Y_{u'm_1} - \mu - \gamma_m[0] \\ Y_{u'm_2} - \mu - \gamma_m[0] \\ \vdots \\ Y_{u'm_g} - \mu - \gamma_m[0] \end{bmatrix}, X_{u'} = \begin{bmatrix} 1, \gamma_{m_1}[1:]^T \\ 1, \gamma_{m_2}[1:]^T \\ \vdots \\ 1, \gamma_{m_g}[1:]^T \end{bmatrix}$$

Taking into account the prior on  $\theta_{u'}$ , we get

$$p(\theta_{u'}|Y, \mu, \sigma^2, \Theta_{-u'}, \Gamma) = N(\mu_{u'}, \Sigma_{u'})$$

$$\text{where } \mu_{u'} = \Sigma_{u'} \frac{1}{\sigma^2} X_{u'}^T Y_{u'}$$


$$\Sigma_{u'}^{-1} = \frac{1}{\sigma^2} X_{u'}^T X_{u'} + \Sigma_{\theta}^{-1}$$

We have  $(U + M) \times (L + 1) + 2$  parameters to estimate. We build a Gibbs sampler to estimate these parameters.

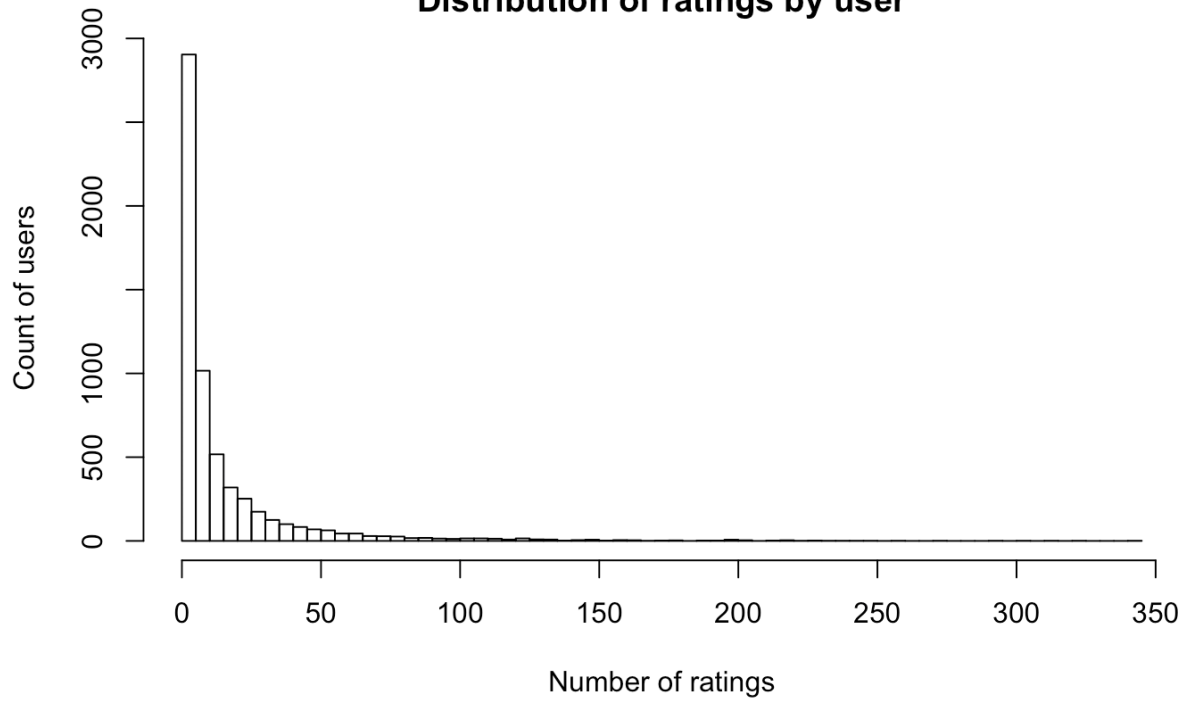
## Data Exploration


We do a basic data exploration and see the distribution of number of ratings by a user, number of ratings for a movie and the rating distribution itself.

**Distribution of ratings for movie**



**Distribution of ratings by user**





To evaluate the results of our model, we divide our data into training, validation and testing data. About 70% is kept as training while the rest is equally divided in validation and testing. While sampling data for training, we made sure that all movies and users that appear in test or validation set have at least two entries in the training data.

## Gibbs Sampling

We now set up Gibbs sampling to evaluate parameters based on our training data. The following values were chosen as the initial values:


- $\mu$ : mean of all ratings in training data
- $\sigma^2$ : variance of all ratings in training data
- $\gamma_m$ :  $\gamma_m[0] = \mu_m - \mu$  where  $\mu_m$  is the mean ratings of movie  $m$ .  $\gamma_m[1:] = 0$
- $\theta_u$ :  $\theta_u[0] = \mu_u - \mu$  where  $\mu_u$  is the mean ratings by user  $u$ .  $\theta_u[1:] = 0$

After removing the burn in samples, the average across all remaining samples was taken to get an estimate of the parameters. These parameters were then used to predict the ratings in validation set. The graphs and details of different trials and their results are included in the Appendix

## Results

Case	L	Samples	Burn in	Training error (RMSE)	Validation error (RMSE)
1	4	500	200	0.66	0.87
2	10	1000	200	0.58	0.88
3	20	1500	1000	0.49	0.89
4	2	5000	4000	0.71	0.88

Based on the results we go with model 4. While the error on validation set is not the least, the value of  $\mu$  had converged for model 4. Running it on the test set we get the following results:



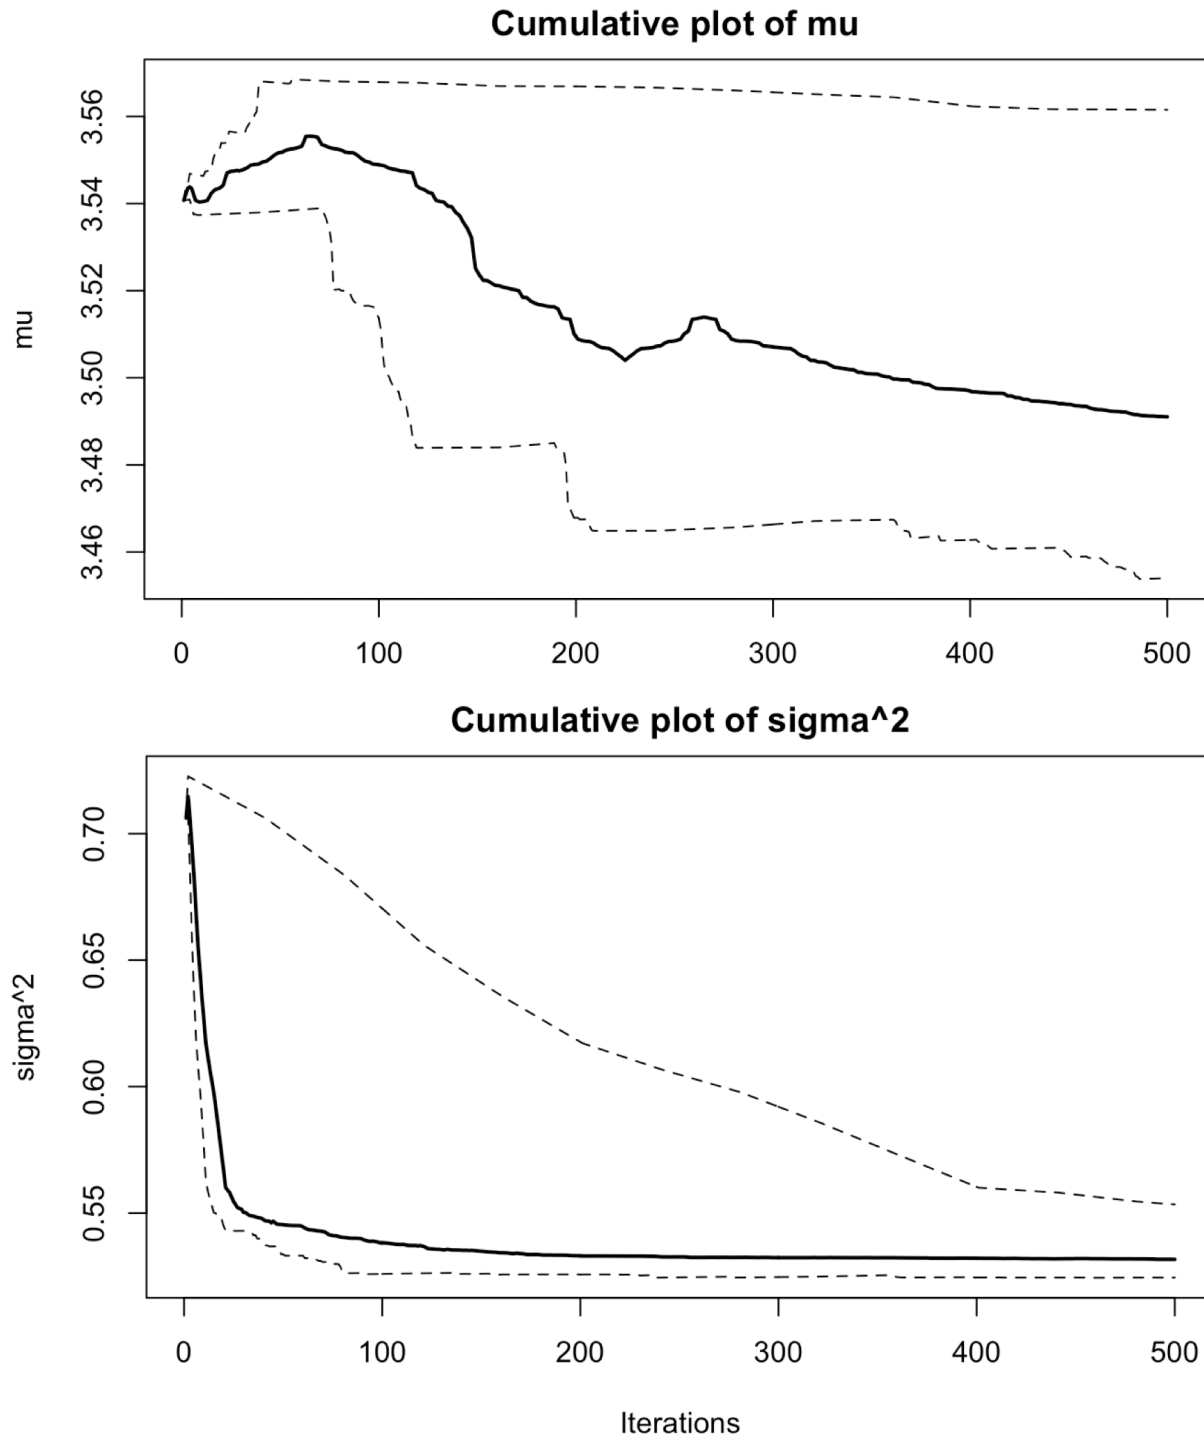
The RMSE on test data was 0.86. The values of  $\mu$  and  $\sigma^2$  were 3.46 and 0.6 respectively.

We see that  $\mu$  does not converge for lower sample sizes. Plotting the predicted data of validation and test set suggests presence of a bias. Running the model for more samples with a higher value of L could result in better accuracy.



# Appendix 1

## Case 1

$L = 4$ , number of samples = 500, burn in = 200




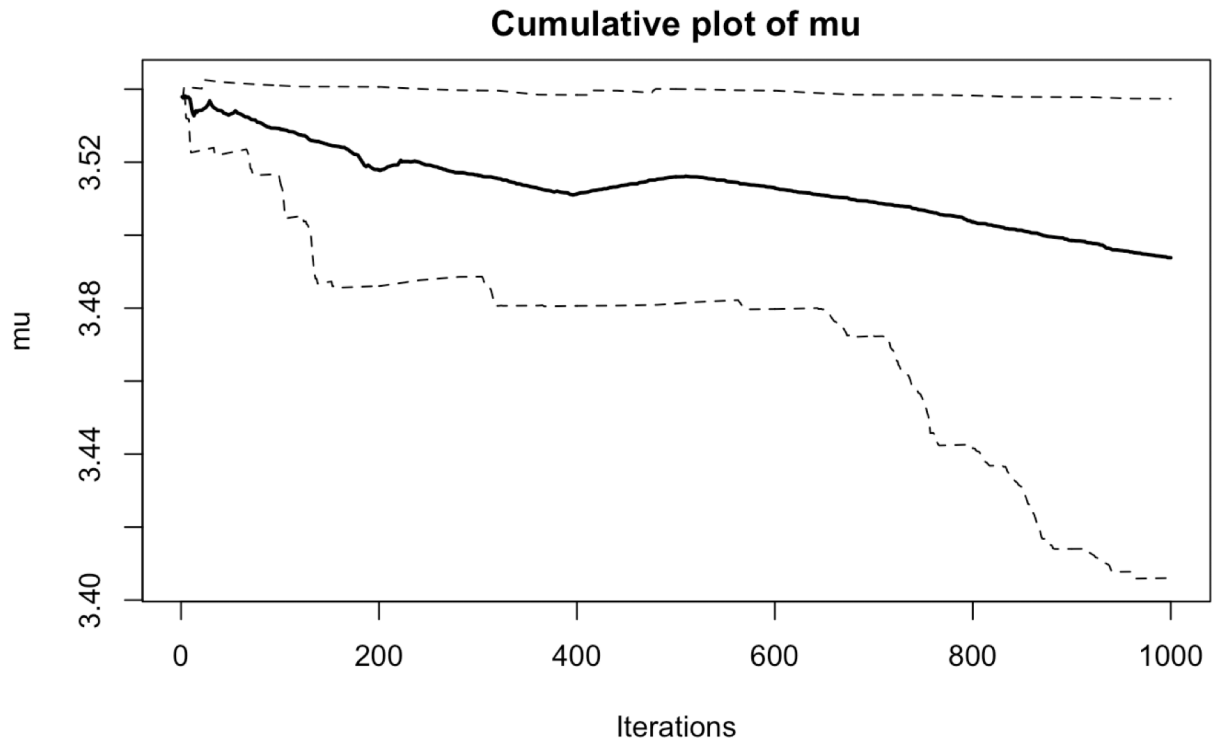





The solid line is the man of predicted ratings and the dotted lines show the range between one standard deviation from the mean

## Case 2


$L = 10$ , number of samples = 1000, burn in = 200



**Comparison on training data**




**Comparison on test data**




### Case 3


$L = 20$ , number of samples = 1500, burn in = 1000



**Comparison on training data**

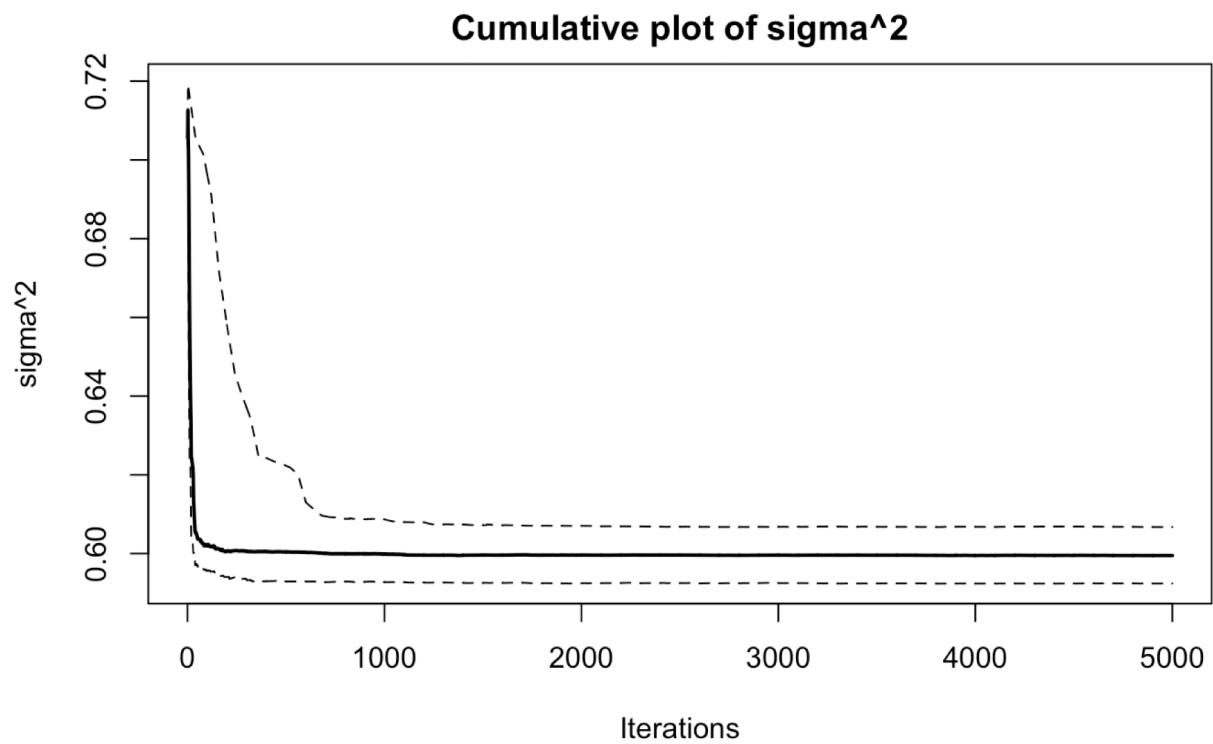



**Comparison on test data**




## Case 4


$L = 2$ , number of samples = 5000, burn in = 4000



**Comparison on training data**



**Comparison on test data**



## Appendix 2

```

1 rm(list = ls())
2 library(invgamma)
3 library(data.table)
4 library(mnormt)
5 library(coda)
6 library(ggplot2)
7
8 gamma_m_sample=function(xm,ym,sigma2,sig_gamma)
9 {
10   inverse=solve(sig_gamma)
11   variance=solve((t(xm)%*%xm)/sigma2+inverse)
12   mean=variance%*%t(xm)%*%ym/sigma2
13   sample=rmnorm(1,mean,variance)
14   return(sample)
15 }
16
17 theta_u_sample=function(xu,yu,sigma2,sig_theta)
18 {
19   inverse=solve(sig_theta)
20   variance=solve((t(xu)%*%xu)/sigma2+inverse)
21   mean=variance%*%t(xu)%*%yu/sigma2
22   sample=rmnorm(1,mean,variance)
23   return(sample)
24 }
25
26 gibbs=function(mdata,L,m,gamma_var,theta_var)
27 {
28   data=copy(mdata)
29   n=length(data$rating)
30
31   total_avg=mean(data$rating)
32   user_avg=aggregate(rating~userId,data,mean)
33   movie_avg=aggregate(rating~movieId,data,mean)
34
35   theta=matrix(0,nrow=length(user_avg[,1]),ncol=L+1)
36   theta[,1]<-user_avg[,2]-total_avg
37   thetasamples=array(0,c(m,length(user_avg[,1]),L+1))
38
39   gamma=matrix(0,nrow=length(movie_avg[,1]),ncol=L+1)
40   gamma[,1]<-movie_avg[,2]-total_avg
41   gammasamples=array(0,c(m,length(movie_avg[,1]),L+1))
42
43   mu=total_avg
44   musamples=array(0,m)
45
46   sigma2=var(data$rating)
47   sigma2samples=array(0,m)
48
49   I=diag(L+1)
50   sig_gamma=I*gamma_var
51   sig_theta=I*theta_var

```



```

52
53 for (i in 1:m)
54 {
55   print(i)
56   y_mu=rowSums(theta[data$uindex,-1]*gamma[data$mindex,-1])+theta[data$
uindex,1]+gamma[data$mindex,1]
57   mu=rnorm(1,mean(data$rating-y_mu),sqrt(sigma2/n))
58   sigma2=rinvgamma(1,n/2,sum((data$rating-y_mu-mu)^2)/2)
59   for (movie in movieids)
60   {
61     index=match(movie,movieids)
62     g_index=which(data$mindex==index)
63     ym=data$rating[g_index]-mu-theta[data$uindex[g_index],1]
64     xm=cbind(1,theta[data$uindex[g_index],-1])
65     gamma[index,]<-gamma_m_sample(xm,ym,sigma2,sig_gamma)
66   }
67   for (user in usersids)
68   {
69     index=match(user,usersids)
70     g_index=which(data$uindex==index)
71     ym=data$rating[g_index]-mu-gamma[data$mindex[g_index],1]
72     xm=cbind(1,gamma[data$mindex[g_index],-1])
73     theta[index,]<-theta_u_sample(xm,ym,sigma2,sig_theta)
74   }
75   musamples[i]<-mu
76   sigma2samples[i]<-sigma2
77   thetasamples[i,]<-theta
78   gammasamples[i,]<-gamma
79 }
80 r=list(codepurple"codepurplemu"codepurple"=musamples,codepurple"codepurplesigma2
codepurple"=sigma2samples,codepurple"codepurpletheta"codepurple"=thetasamples,
codepurple"codepurplegamma"codepurple"=gammasamples)
81 return(r)
82 }
83
84 mcmcplots=function(gibbsresult)
85 {
86   cumuplot(mcmc(gibbsresult$mu),main=codepurple"Cumulative"codepurple
codepurpleplotcodepurple codepurpleofcodepurple codepurplemu"codepurple",ylab=codepurple"
codepurplemu"codepurple")
87   cumuplot(mcmc(gibbsresult$sigma2),main=codepurple"Cumulative"codepurple
codepurpleplotcodepurple codepurpleofcodepurple codepurplesigmacodepurple^2codepurple",
ylab=codepurple"codepurplesigmacodepurple^2codepurple")
88   cumuplot(mcmc(gibbsresult$gamma[,1,]),)
89   cumuplot(mcmc(gibbsresult$theta[,1,]),)
90   plot(mcmc(gibbsresult$mu),trace=F)
91   plot(mcmc(gibbsresult$sigma2),trace=F)
92   plot(mcmc(gibbsresult$gamma[,1,]),trace=F)
93   plot(mcmc(gibbsresult$theta[,1,]),trace=F)
94   print(effectiveSize(gibbsresult$mu))
95   print(effectiveSize(gibbsresult$sigma2))
96   print(effectiveSize(gibbsresult$gamma[,1,]))
97   print(effectiveSize(gibbsresult$theta[,1,]))
98 }

```

```

99
100 plot_results=function(data, gibbsresult, burnin)
101 {
102   m=length(gibbsresult$mu)
103   mu_final=mean(gibbsresult$mu[c(burnin:m)])
104   sigma2final=mean(gibbsresult$sigma2[c(burnin:m)])
105   gamma_final=colMeans(gibbsresult$gamma[c(burnin:m),,], dim=1)
106   theta_final=colMeans(gibbsresult$theta[c(burnin:m),,], dim=1)
107
108   predicted=mu_final+gamma_final[data$minindex,1]+theta_final[data$uindex,1] +
     rowSums(theta_final[data$uindex,-1]*gamma_final[data$minindex,-1])
109
110   compare=as.data.frame(predicted)
111   compare$actual=data$rating
112
113   hist(compare$actual-compare$predicted, main="codepurple" codepurpleDistribution
     codepurple codepurpleofcodepurple codepurpleerrorscodepurple codepurpleincodepurple
     codepurpletrainingcodepurple codepurpledatacodepurple", xlab="codepurple" codepurpleError
     codepurple", ylab="codepurple" codepurpleFrequencycodepurple")
114   print(sd(compare$actual-compare$predicted))
115
116   meantable=aggregate(predicted~actual, compare, mean)
117   sdtable=aggregate(predicted~actual, compare, sd)
118
119   plot(compare$actual, compare$predicted, main="codepurple" codepurpleComparison
     codepurple codepurpleoncodepurple codepurpletrainingcodepurple codepurpledatacodepurple",
     xlab="codepurple" codepurpleActualcodepurple codepurpleratingcodepurple", ylab="codepurple"
     codepurplePredictedcodepurple codepurpleratingcodepurple")
120   lines(meantable$actual, meantable$predicted, lwd=2)
121   lines(sdtable$actual, meantable$predicted+sdtable$predicted, lty=2)
122   lines(sdtable$actual, meantable$predicted-sdtable$predicted, lty=2)
123 }
124
125 validation_results=function(gibbsresult, testdata, burnin)
126 {
127   m=length(gibbsresult$mu)
128   mu_final=mean(gibbsresult$mu[c(burnin:m)])
129   sigma2final=mean(gibbsresult$sigma2[c(burnin:m)])
130   gamma_final=colMeans(gibbsresult$gamma[c(burnin:m),,], dim=1)
131   theta_final=colMeans(gibbsresult$theta[c(burnin:m),,], dim=1)
132
133   predicted=mu_final+gamma_final[testdata$minindex,1]+theta_final[testdata$
     uindex,1] + rowSums(theta_final[testdata$uindex,-1]*gamma_final[testdata$
     minindex,-1])
134
135   compare=as.data.frame(predicted)
136   compare$actual=testdata$rating
137
138   hist(compare$actual-compare$predicted, main="codepurple" codepurpleDistribution
     codepurple codepurpleofcodepurple codepurpleerrorscodepurple codepurpleincodepurple
     codepurpletestcodepurple codepurpledatacodepurple", xlab="codepurple" codepurpleError
     codepurple", ylab="codepurple" codepurpleFrequencycodepurple")
139   print(sd(compare$actual-compare$predicted))
140

```

```

141  meantable=aggregate(predicted~actual,compare,mean)
142  sdtable=aggregate(predicted~actual,compare,sd)
143
144  plot(compare$actual,compare$predicted,main=codepurple"codepurpleComparison
      codepurple codepurpleoncodepurple codepurletestcodepurple codepurpledatacodepurple",xlab
      =codepurple"codepurpleActualcodepurple codepurpleratingcodepurple",ylab=codepurple"
      codepurplePredictedcodepurple codepurpleratingcodepurple")
145  lines(meantable$actual,meantable$predicted,lwd=2)
146  lines(sdtable$actual,meantable$predicted+sdtable$predicted,lty=2)
147  lines(sdtable$actual,meantable$predicted-sdtable$predicted,lty=2)
148 }
149
150 movie=read.csv(codepurple"codepurpleratings_codepurplesmallcodepurple.codepurplecsv
      codepurple")
151 perusercount=aggregate(rating~userId,movie,length)
152 hist(perusercount$rating,breaks=50,main=codepurple"codepurpleDistributioncodepurple
      codepurpleofcodepurple codepurpleratingscodepurple codepurpleforcodepurple
      codepurplemoviecodepurple", xlab=codepurple"codepurpleNumbercodepurple codepurpleof
      codepurple codepurpleratingscodepurple",ylab=codepurple"codepurpleCountcodepurple
      codepurpleofcodepurple codepurplemoviescodepurple")
153 permoviecount=aggregate(rating~movieId,movie,length)
154 hist(permoviecount$rating,breaks=50,main=codepurple"codepurpleDistribution
      codepurple codepurpleofcodepurple codepurpleratingscodepurple codepurplebycodepurple
      codepurpleusercodepurple", xlab=codepurple"codepurpleNumbercodepurple codepurpleof
      codepurple codepurpleratingscodepurple",ylab=codepurple"codepurpleCountcodepurple
      codepurpleofcodepurple codepurpleuserscodepurple")
155 hist(movie$rating,breaks=5,main=codepurple"codepurpleDistributioncodepurple
      codepurpleofcodepurple codepurpleratingscodepurple", xlab=codepurple"codepurpleRating
      codepurple",ylab=codepurple"codepurpleCountcodepurple codepurpleofcodepurple
      codepurplemoviescodepurple")
156
157 n=length(movie$rating)
158
159 total_avg=mean(movie$rating)
160 user_avg=aggregate(rating~userId,movie,mean)
161 movie_avg=aggregate(rating~movieId,movie,mean)
162
163 userids=user_avg[,1]
164 movieids=movie_avg[,1]
165
166 userindex=match(movie$userId,userids)
167 movieindex=match(movie$movieId,movieids)
168
169 movie$uindex<-userindex
170 movie$mindex<-movieindex
171
172 test_ind <- sample(seq_len(n), size = floor(0.32*n))
173
174 temp=as.data.frame(table(movie$movieId[-test_ind]))
175 allowed=as.data.frame(temp[temp$Freq>2,]$Var1)
176 names(allowed)<-(codepurple"codepurplemoviecodepurple")
177 test_ind=test_ind[movie$movieId[test_ind]%in%allowed$movie]
178
179 temp=as.data.frame(table(movie$userId[-test_ind]))

```

```

180 allowed=as.data.frame(temp[temp$Freq>2,]$Var1)
181 names(allowed)<-(codepurple"codepurpleusercodepurple")
182 test_ind=test_ind[movie$userId[test_ind]%in%allowed$user]
183
184 movie_train=movie[-test_ind,]
185
186 validation_ind=sample(test_ind, size = floor(length(test_ind)/2))
187 test_ind=test_ind[!test_ind%in%validation_ind]
188
189 movie_validate=movie[validation_ind,]
190 movie_test=movie[test_ind,]
191
192 test2=gibbs(movie_train,4,500,1,1)
193 mcmcplots(test2)
194 plot_results(movie_train,test2,200)
195 validation_results(test2,movie_test,200)
196
197 test3=gibbs(movie_train,10,1000,1,1)
198 mcmcplots(test3)
199 plot_results(movie_train,test3,200)
200 validation_results(test3,movie_validate,200)
201
202 test4=gibbs(movie_train,20,1500,1,1)
203 mcmcplots(test4)
204 plot_results(movie_train,test4,1000)
205 validation_results(test4,movie_validate,1000)
206
207 test5=gibbs(movie_train,2,5000,1,1)
208 mcmcplots(test5)
209 plot_results(movie_train,test5,4000)
210 validation_results(test5,movie_validate,4000)
211 validation_results(test5,movie_test,4000)
212 print(mean(test5$mu[c(4000:5000)]))
213 print(mean(test5$sigma2[c(4000:5000)]))

```